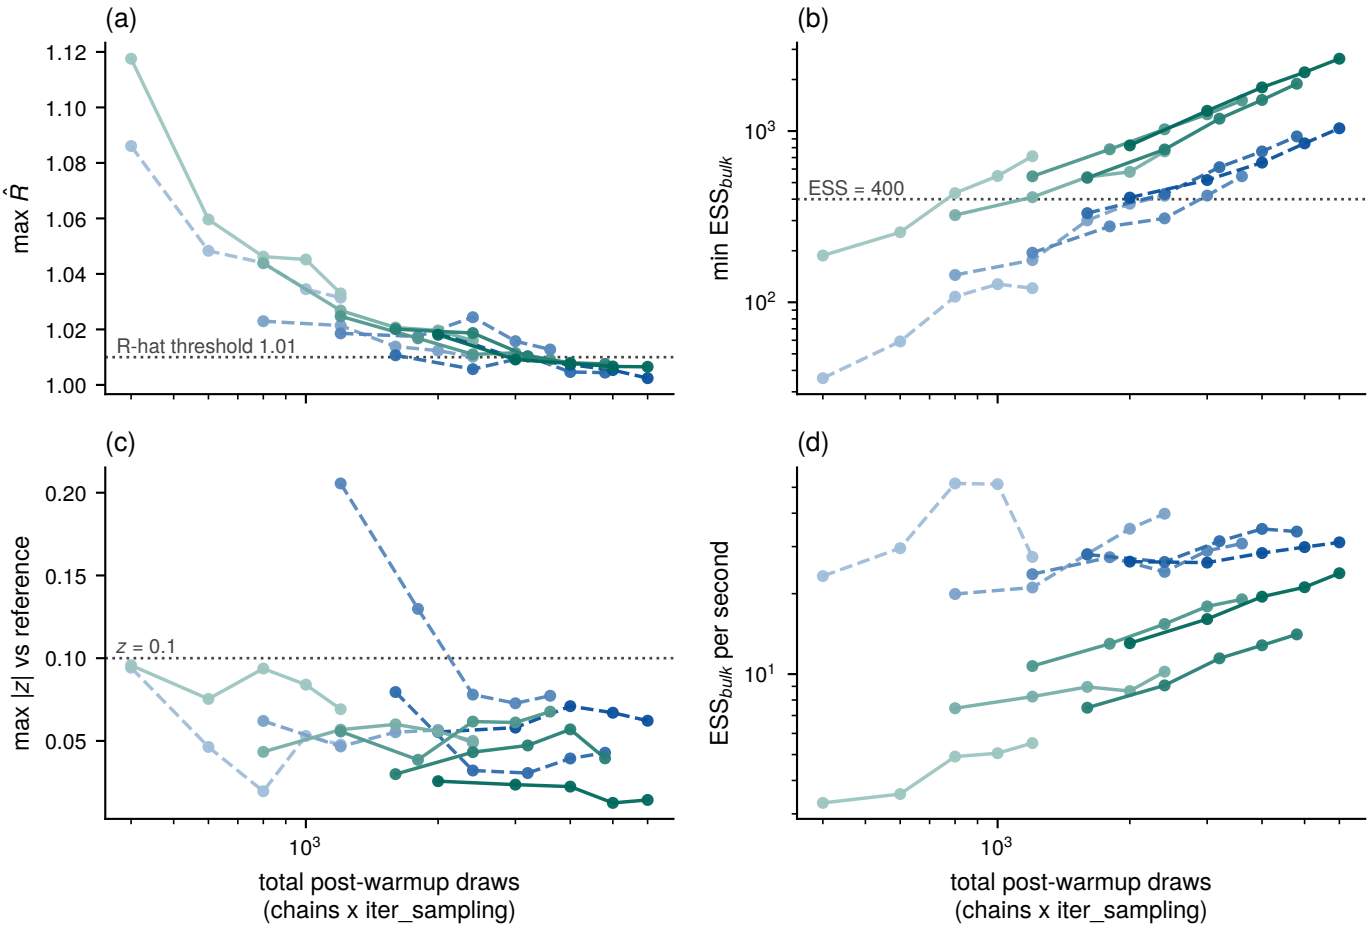
Convergence, accuracy and efficiency vs sampling budget
(a)
max $\hat{R}$
1.12
1.10
1.08
1.06
1.04
1.02
1.00
R-hat threshold 1.01
(b)
min ESS_{bulk}
10^3
10^2
ESS = 400
(c)
max |z| vs reference
0.20
0.15
0.10
0.05
z = 0.1
10^3
total post-warmup draws
(chains x iter_sampling)
(d)
ESS_{bulk} per second
10^1
10^3
total post-warmup draws
(chains x iter_sampling)
univariate (prior-approximation)
EIV multivariate, T_0-shift