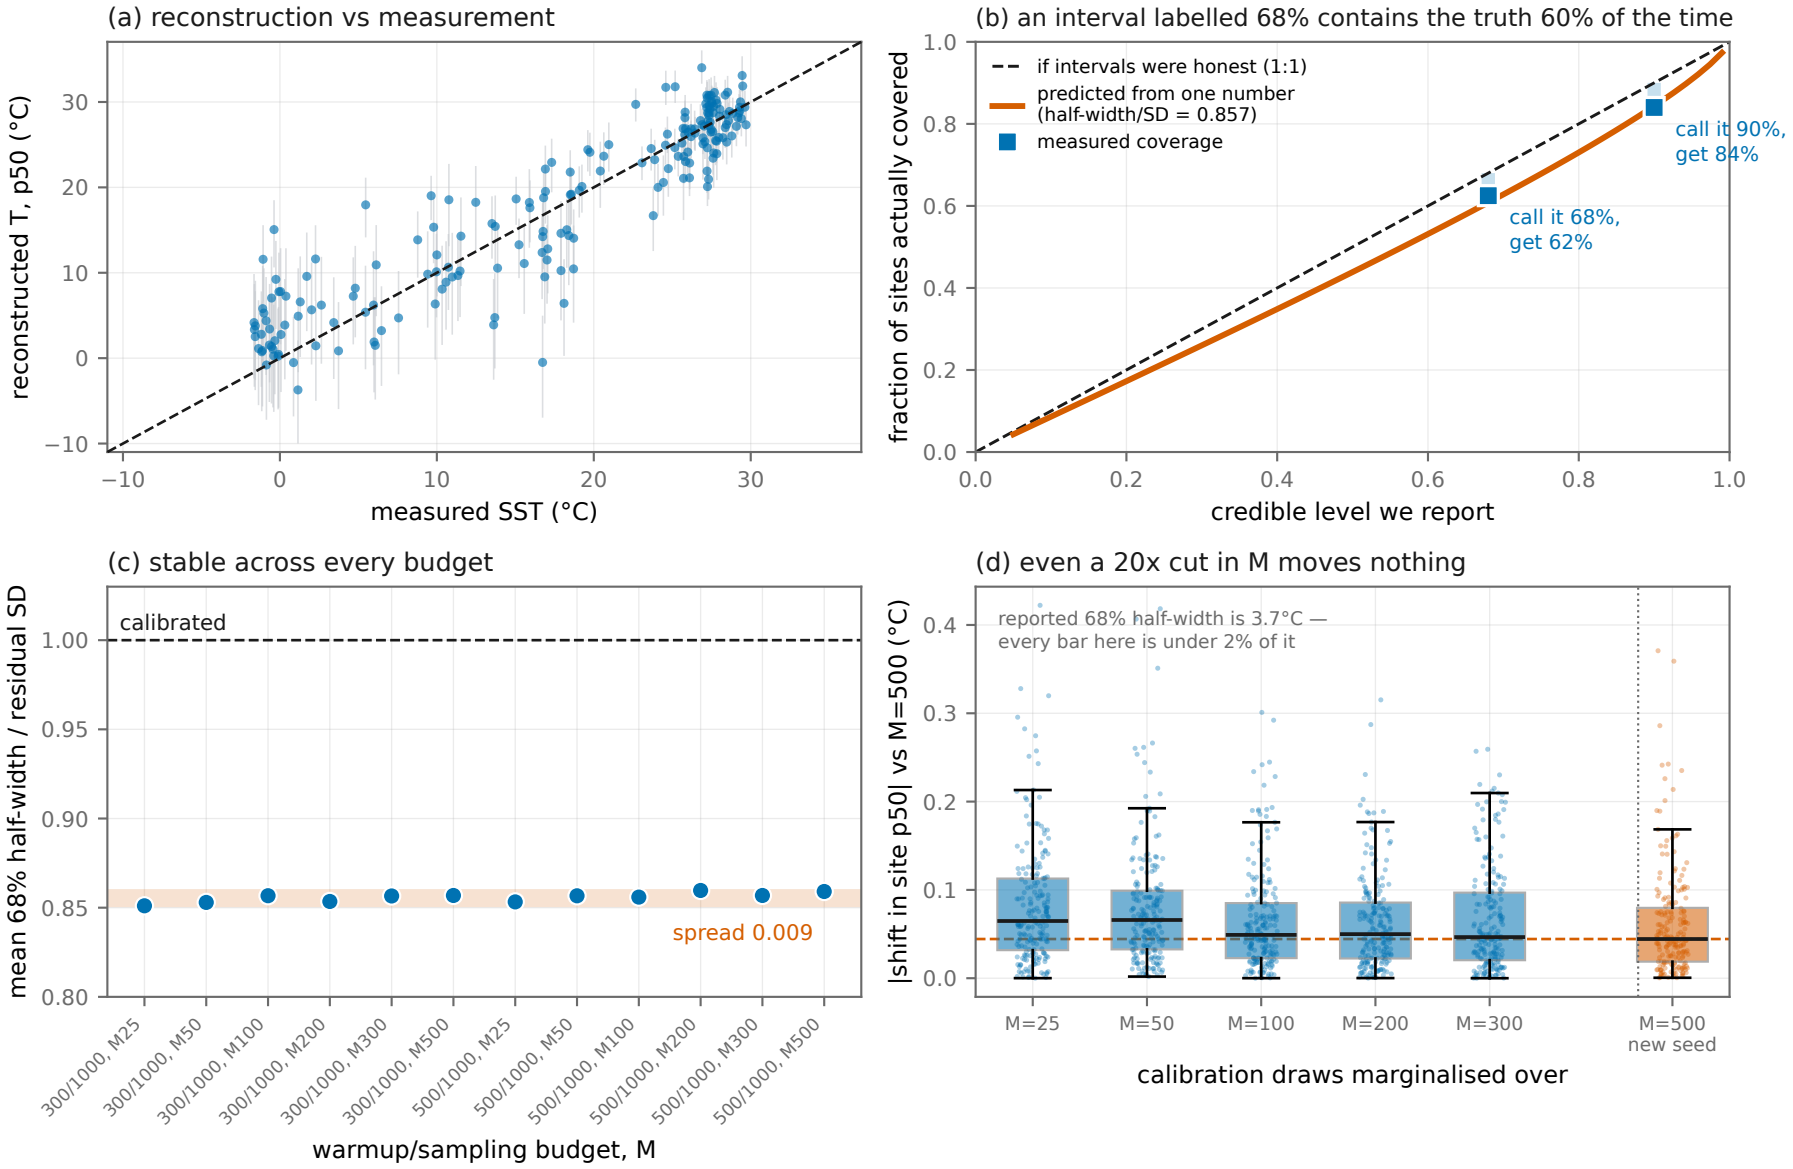
Predictive-interval calibration of the inverse model (n = 200 coretop sites)
(a) reconstruction vs measurement
reconstructed T, p50 (°C)
measured SST (°C)
(b) an interval labelled 68% contains the truth 60% of the time
fraction of sites actually covered
credible level we report
if intervals were honest (1:1)
predicted from one number
(half-width/SD = 0.857)
measured coverage
call it 68%,
get 62%
call it 90%,
get 84%
(c) stable across every budget
mean 68% half-width / residual SD
calibrated
spread 0.009
warmup/sampling budget, M
300/1000, M25
300/1000, M50
300/1000, M100
300/1000, M200
300/1000, M300
300/1000, M500
500/1000, M25
500/1000, M50
500/1000, M100
500/1000, M200
500/1000, M300
500/1000, M500
(d) even a 20x cut in M moves nothing
|shift in site p50| vs M=500 (°C)
reported 68% half-width is 3.7°C —
every bar here is under 2% of it
M=25
M=50
M=100
M=200
M=300
M=500
new seed
calibration draws marginalised over